

Class09_mini_project

Gabrielle Meza (A13747395)

10/27/2021

R Markdown

Inputing data, and setting up names. we dont want diagnosis becuase that is what we are trying to find, so we ar egetting rid of it in our dataset and then storing it as another piece for later refernece

```
fna.data <- "WisconsinCancer.csv"
wisc.df <- read.csv(fna.data, row.names = 1)
head(wisc.df)
```

```
##      diagnosis radius_mean texture_mean perimeter_mean area_mean
## 842302      M      17.99      10.38      122.80      1001.0
## 842517      M      20.57      17.77      132.90      1326.0
## 84300903     M      19.69      21.25      130.00      1203.0
## 84348301     M      11.42      20.38       77.58       386.1
## 84358402     M      20.29      14.34      135.10      1297.0
## 843786      M      12.45      15.70       82.57       477.1
##      smoothness_mean compactness_mean concavity_mean concave.points_mean
## 842302      0.11840      0.27760      0.3001      0.14710
## 842517      0.08474      0.07864      0.0869      0.07017
## 84300903     0.10960      0.15990      0.1974      0.12790
## 84348301     0.14250      0.28390      0.2414      0.10520
## 84358402     0.10030      0.13280      0.1980      0.10430
## 843786      0.12780      0.17000      0.1578      0.08089
##      symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
## 842302      0.2419      0.07871      1.0950      0.9053      8.589
## 842517      0.1812      0.05667      0.5435      0.7339      3.398
## 84300903     0.2069      0.05999      0.7456      0.7869      4.585
## 84348301     0.2597      0.09744      0.4956      1.1560      3.445
## 84358402     0.1809      0.05883      0.7572      0.7813      5.438
## 843786      0.2087      0.07613      0.3345      0.8902      2.217
##      area_se smoothness_se compactness_se concavity_se concave.points_se
## 842302     153.40      0.006399      0.04904      0.05373      0.01587
## 842517      74.08      0.005225      0.01308      0.01860      0.01340
## 84300903     94.03      0.006150      0.04006      0.03832      0.02058
## 84348301     27.23      0.009110      0.07458      0.05661      0.01867
## 84358402     94.44      0.011490      0.02461      0.05688      0.01885
## 843786      27.19      0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
## 842302      0.03003      0.006193      25.38      17.33
## 842517      0.01389      0.003532      24.99      23.41
## 84300903     0.02250      0.004571      23.57      25.53
```

```
wisc.data <- wisc.df[,-1]
diagnosis<- as.factor(wisc.df$diagnosis)
diagnosis
```

```
is.vector(diagnosis)
```

```
nrow(wisc.data)
```

```
## [1] 569
```

```
dim(wisc.data)
```

```
## [1] 569 30
```

Q2. How many of the observations have a malignant diagnosis?

= assigns a value, == pulls to find. also table(diagnosis)works

```
sum(diagnosis=="M")
```

```
## [1] 212
```

```
sum(diagnosis=="B")
```

```
## [1] 357
```

```
table(diagnosis)
```

```
## diagnosis  
##    B    M  
## 357 212
```

Q3. How many variables/features in the data are suffixed with _mean?

```
grep("_mean", colnames(wisc.data), value= TRUE )
```

```
## [1] "radius_mean"      "texture_mean"      "perimeter_mean"  
## [4] "area_mean"        "smoothness_mean"   "compactness_mean"  
## [7] "concavity_mean"    "concave.points_mean" "symmetry_mean"  
## [10] "fractal_dimension_mean"
```

```
meanvar <- grep("_mean", colnames(wisc.data), value= TRUE )  
length(meanvar)
```

```
## [1] 10
```

#Check your data scaling is correct

```
colMeans(wisc.data)
```

```
##          radius_mean      texture_mean      perimeter_mean
##      1.412729e+01      1.928965e+01      9.196903e+01
##          area_mean      smoothness_mean      compactness_mean
##      6.548891e+02      9.636028e-02      1.043410e-01
##      concavity_mean      concave.points_mean      symmetry_mean
##      8.879932e-02      4.891915e-02      1.811619e-01
## fractal_dimension_mean      radius_se      texture_se
##      6.279761e-02      4.051721e-01      1.216853e+00
##      perimeter_se      area_se      smoothness_se
##      2.866059e+00      4.033708e+01      7.040979e-03
##      compactness_se      concavity_se      concave.points_se
##      2.547814e-02      3.189372e-02      1.179614e-02
##      symmetry_se      fractal_dimension_se      radius_worst
##      2.054230e-02      3.794904e-03      1.626919e+01
##      texture_worst      perimeter_worst      area_worst
##      2.567722e+01      1.072612e+02      8.805831e+02
##      smoothness_worst      compactness_worst      concavity_worst
##      1.323686e-01      2.542650e-01      2.721885e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      1.146062e-01      2.900756e-01      8.394582e-02
```

```
apply(wisc.data,2,sd)
```

```
##          radius_mean      texture_mean      perimeter_mean
##      3.524049e+00      4.301036e+00      2.429898e+01
##          area_mean      smoothness_mean      compactness_mean
##      3.519141e+02      1.406413e-02      5.281276e-02
##      concavity_mean      concave.points_mean      symmetry_mean
##      7.971981e-02      3.880284e-02      2.741428e-02
## fractal_dimension_mean      radius_se      texture_se
##      7.060363e-03      2.773127e-01      5.516484e-01
##      perimeter_se      area_se      smoothness_se
##      2.021855e+00      4.549101e+01      3.002518e-03
##      compactness_se      concavity_se      concave.points_se
##      1.790818e-02      3.018606e-02      6.170285e-03
##      symmetry_se      fractal_dimension_se      radius_worst
##      8.266372e-03      2.646071e-03      4.833242e+00
##      texture_worst      perimeter_worst      area_worst
##      6.146258e+00      3.360254e+01      5.693570e+02
##      smoothness_worst      compactness_worst      concavity_worst
##      2.283243e-02      1.573365e-01      2.086243e-01
##      concave.points_worst      symmetry_worst      fractal_dimension_worst
##      6.573234e-02      6.186747e-02      1.806127e-02
```

```
apply
```

```
## function (X, MARGIN, FUN, ..., simplify = TRUE)
## {
##     FUN <- match.fun(FUN)
##     simplify <- isTRUE(simplify)
##     dl <- length(dim(X))
##     if (!dl)
##         stop("dim(X) must have a positive length")
```

```

##   if (is.object(X))
##       X <- if (d1 == 2L)
##           as.matrix(X)
##       else as.array(X)
##   d <- dim(X)
##   dn <- dimnames(X)
##   ds <- seq_len(d1)
##   if (is.character(MARGIN)) {
##       if (is.null(dnn <- names(dn)))
##           stop("'X' must have named dimnames")
##       MARGIN <- match(MARGIN, dnn)
##       if (anyNA(MARGIN))
##           stop("not all elements of 'MARGIN' are names of dimensions")
##   }
##   d.call <- d[-MARGIN]
##   d.ans <- d[MARGIN]
##   if (anyNA(d.call) || anyNA(d.ans))
##       stop("'MARGIN' does not match dim(X)")
##   s.call <- ds[-MARGIN]
##   s.ans <- ds[MARGIN]
##   dn.call <- dn[-MARGIN]
##   dn.ans <- dn[MARGIN]
##   d2 <- prod(d.ans)
##   if (d2 == 0L) {
##       newX <- array(vector(typeof(X), 1L), dim = c(prod(d.call),
##           1L))
##       ans <- forceAndCall(1, FUN, if (length(d.call) < 2L) newX[,
##           1] else array(newX[, 1L], d.call, dn.call), ...)
##       return(if (is.null(ans)) ans else if (length(d.ans) <
##           2L) ans[1L][-1L] else array(ans, d.ans, dn.ans))
##   }
##   newX <- aperm(X, c(s.call, s.ans))
##   dim(newX) <- c(prod(d.call), d2)
##   ans <- vector("list", d2)
##   if (length(d.call) < 2L) {
##       if (length(dn.call))
##           dimnames(newX) <- c(dn.call, list(NULL))
##       for (i in 1L:d2) {
##           tmp <- forceAndCall(1, FUN, newX[, i], ...)
##           if (!is.null(tmp))
##               ans[[i]] <- tmp
##       }
##   }
##   else for (i in 1L:d2) {
##       tmp <- forceAndCall(1, FUN, array(newX[, i], d.call,
##           dn.call), ...)
##       if (!is.null(tmp))
##           ans[[i]] <- tmp
##   }
##   ans.list <- !simplify || is.recursive(ans[[1L]])
##   l.ans <- length(ans[[1L]])
##   ans.names <- names(ans[[1L]])
##   if (!ans.list)
##       ans.list <- any(lengths(ans) != l.ans)

```

```

##   if (!ans.list && length(ans.names)) {
##       all.same <- vapply(ans, function(x) identical(names(x),
##           ans.names), NA)
##       if (!all(all.same))
##           ans.names <- NULL
##   }
##   len.a <- if (ans.list)
##       d2
##   else length(ans <- unlist(ans, recursive = FALSE))
##   if (length(MARGIN) == 1L && len.a == d2) {
##       names(ans) <- if (length(dn.ans[[1L]]))
##           dn.ans[[1L]]
##       ans
##   }
##   else if (len.a == d2)
##       array(ans, d.ans, dn.ans)
##   else if (len.a && len.a%%d2 == 0L) {
##       if (is.null(dn.ans))
##           dn.ans <- vector(mode = "list", length(d.ans))
##       dn1 <- list(ans.names)
##       if (length(dn.call) && !is.null(n1 <- names(dn <- dn.call[1])) &&
##           nzchar(n1) && length(ans.names) == length(dn[[1]]))
##           names(dn1) <- n1
##       dn.ans <- c(dn1, dn.ans)
##       array(ans, c(len.a%%d2, d.ans), if (!is.null(names(dn.ans)) ||
##           !all(vapply(dn.ans, is.null, NA)))
##           dn.ans)
##   }
##   else ans
## }
## <bytecode: 0x7fcd39dd7580>
## <environment: namespace:base>

```

```

wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)

```

```

## Importance of components:
##
##      PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##
##      PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##
##      PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##
##      PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##
##      PC29     PC30

```

```
## Standard deviation      0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion  1.00000 1.00000
```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27%


Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

Up to PC3

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

up to PC7


```
biplot(wisc.pr)
```



Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

It is super jumbled. You can not really tell anything from this, it is so clumped together.


```
plot(wisc.pr$x[,1:2], col=diagnosis)
```




Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

In the first plot, they are more differentiated, and spread out. Because principal component 2 explains more variance in the original data than principal component 3, you can see that the first plot has a cleaner cut separating the two subgroups

```
plot(wisc.pr$x[,1], wisc.pr$x[,3], col=diagnosis, xlab="PC1", ylab="PC3")
```

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis
library(ggplot2)
ggplot(df) +
  aes(PC1, PC2, col= diagnosis) +
  geom_point()
```

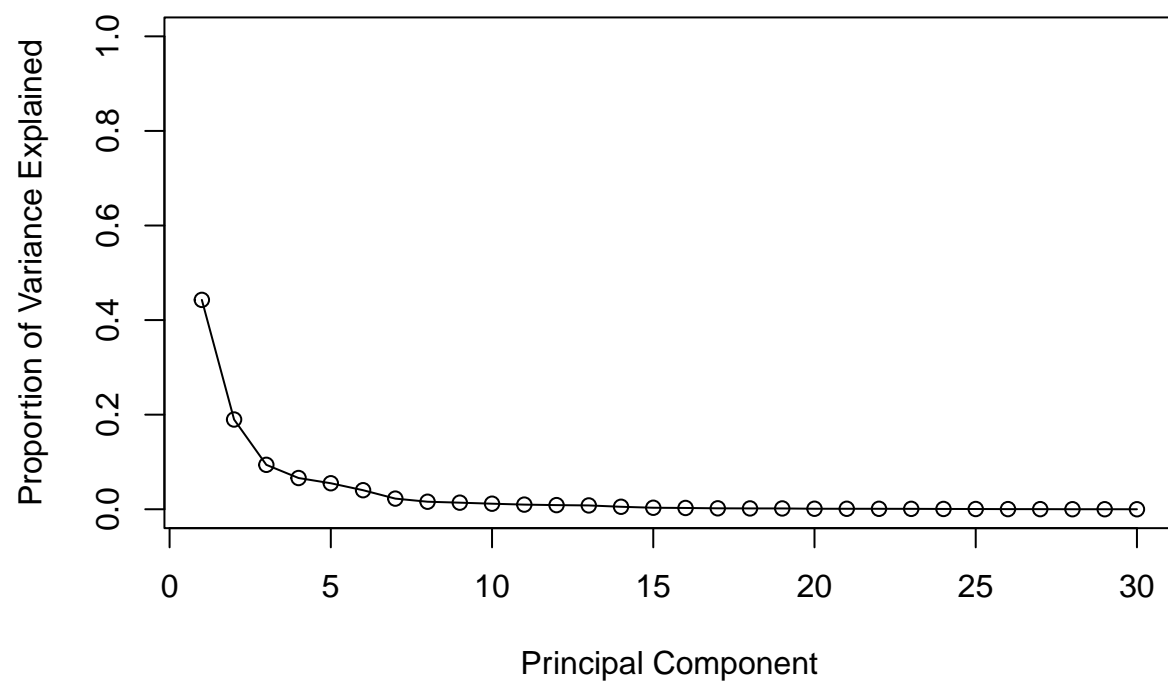


```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```


```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

```
pve <- pr.var/ sum(pr.var)
```

```
plot(pve, xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



```
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

-0.26085376. this is the amount of shift in the axis rotation

```
wisc.pr$rotation[,1]
```

```
##          radius_mean          texture_mean          perimeter_mean
##          -0.21890244          -0.10372458          -0.22753729
##          area_mean          smoothness_mean          compactness_mean
##          -0.22099499          -0.14258969          -0.23928535
##          concavity_mean          concave.points_mean          symmetry_mean
##          -0.25840048          -0.26085376          -0.13816696
## fractal_dimension_mean          radius_se          texture_se
##          -0.06436335          -0.20597878          -0.01742803
##          perimeter_se          area_se          smoothness_se
##          -0.21132592          -0.20286964          -0.01453145
##          compactness_se          concavity_se          concave.points_se
##          -0.17039345          -0.15358979          -0.18341740
##          symmetry_se          fractal_dimension_se          radius_worst
##          -0.04249842          -0.10256832          -0.22799663
##          texture_worst          perimeter_worst          area_worst
##          -0.10446933          -0.23663968          -0.22487053
##          smoothness_worst          compactness_worst          concavity_worst
##          -0.12795256          -0.21009588          -0.22876753
```

```
## concave.points_worst      symmetry_worst fractal_dimension_worst
##          -0.25088597          -0.12290456          -0.13178394
```

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

You need 4 PCs to describe 80% of the variance. but that is with rounding up, so really 5

```
var <- summary(wisc.pr)
var
```

```
## Importance of components:
```


```
##          PC1    PC2    PC3    PC4    PC5    PC6    PC7
## Standard deviation    3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##          PC8    PC9    PC10   PC11   PC12   PC13   PC14
## Standard deviation    0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##          PC15   PC16   PC17   PC18   PC19   PC20   PC21
## Standard deviation    0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##          PC22   PC23   PC24   PC25   PC26   PC27   PC28
## Standard deviation    0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##          PC29   PC30
## Standard deviation    0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

```
sum(var$importance[3,] < 0.8)
```

```
## [1] 4
```

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist, "complete")
wisc.hclust.other <- hclust(data.dist, "ward.D2")
plot(wisc.hclust.other)
```

Cluster Dendrogram




```
data.dist  
hclust (*, "ward.D2")
```

Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

at height 19

```
plot(wisc.hclust)  
abline(h=19, col="red", lty=2)
```

Cluster Dendrogram



```
data.dist
hclust(*, "complete")
```

```
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##                   1 12 165
##                   2  2  5
##                   3 343 40
##                   4  0  2
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

4 and 5 seem to be good clusters because they separate each diagnosis, and there are not a large amount of clusters with very few diagnosis clouding up the data.

```
wisc.hclust.clusters.test <- cutree(wisc.hclust, k = 5)
table(wisc.hclust.clusters.test, diagnosis)
```


```
##              diagnosis
## wisc.hclust.clusters.test  B  M
##                   1 12 165
##                   2  0  5
##                   3 343 40
##                   4  2  0
##                   5  0  2
```

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

i liked ward.D2 the most because it has a major set of clusters early on that define 2 groups, and then more branching occurs. this matches that in the cluster test there are 2 major groups that define B and M

```
wisc.hclust.other <- hclust(data.dist, "single")  
plot(wisc.hclust.other)
```

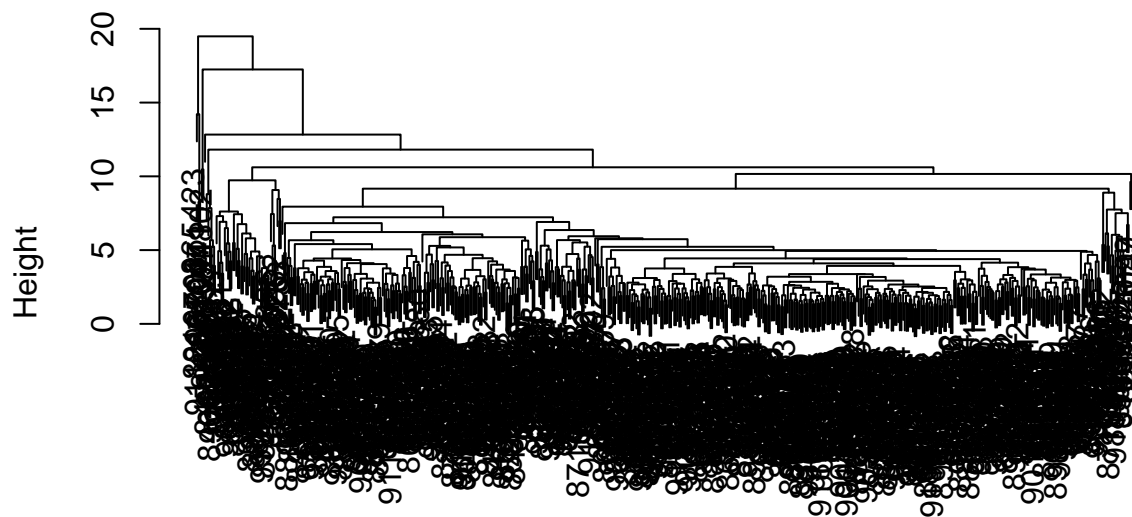
Cluster Dendrogram



data.dist
hclust (*, "single")

```
wisc.hclust.other1 <- hclust(data.dist, "average")  
plot(wisc.hclust.other1)
```



Cluster Dendrogram



data.dist
hclust (*, "average")

```
wisc.hclust.other2 <- hclust(data.dist, "ward.D2")  
plot(wisc.hclust.other2)
```

Cluster Dendrogram




```
data.dist  
hclust (*, "ward.D2")
```

##Combining Methods- Cluster my PCA results

I will use 4 PCS (Covers 80%) and 'hclust()' and 'dist()' as an input. Has to be a distance matrix input

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:4]), method = "ward.D2")  
plot(wisc.pr.hclust)
```

Cluster Dendrogram




```
dist(wisc.pr$x[, 1:4])  
hclust (*, "ward.D2")
```

This cluster is showing me more separation than the other one. This branching pattern is more structured. 2 subgroups can still be seen

```
plot(wisc.pr.hclust)  
abline(h=80, col= "orange")
```

Cluster Dendrogram



```
dist(wisc.pr$x[, 1:4])
hclust (*, "ward.D2")
```

Let's find our cluster membership vector by cutting this tree into $k=2$ groups. We will name this `grps` as just a group separation display of data. Then use `table` to see the groupings.

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
## grps
##   1   2
## 171 398
```

Now let's compare the expert M and B vector

```
table(diagnosis)
```

```
## diagnosis
##   B   M
## 357 212
```

We can do a cross-table by giving the `'table()'` function two inputs.

We want to know how many Ms are in a group and how many Bs are in the other group

```
table(grps, diagnosis)
```

```
##      diagnosis
## grps   B    M
##    1    6 165
##    2 351   47
```

True positives: 165, True Negatives: 351, False Positives: 6, False negatives: 6 **Accuracy**, essentially how many did we get correct?

About 89% accuracy, pretty good.

```
(161+351) / nrow(wisc.data)
```

```
## [1] 0.8998243
```

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

It seems to be pretty good, 89% is good, but not great. Using both methods did make this grouping better.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

‘Did not do kmeans’

They do look similar, but post PCA had more clear separation.

```
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B    M
##                   1  12 165
##                   2   2   5
##                   3 343  40
##                   4   0   2
```

```
table(grps, diagnosis)
```

```
##      diagnosis
## grps   B    M
##    1    6 165
##    2 351   47
```

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity?
How about sensitivity?

This seems to be a more specific analysis. I would prefer the more sensitive analysis, better safe than sorry.

Sensitivity: test ability to correctly detect ill patients with the condition. $(TP/(TP+FN))$

Specificity: test ability to correctly reject healthy patients without a condition. $(TN/(TN+FN))$

```
(165/(165+47))
```

```
## [1] 0.7783019
```

```
(351/(351+47))
```

```
## [1] 0.8819095
```

Prediction

We will use the `predict()` function that will take our PCA model from before and new cancer cell data and project that data onto our PCA space. `npc` is what we are creating


```
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6          PC7
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
##          PC8          PC9          PC10          PC11          PC12          PC13          PC14
## [1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
## [2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
##          PC15          PC16          PC17          PC18          PC19          PC20
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
##          PC21          PC22          PC23          PC24          PC25          PC26
## [1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##          PC27          PC28          PC29          PC30
## [1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
## [2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

Now we will add these samples to our PCA plot.

`pch` is the shape `cex` is the size

```
plot(wisc.pr$x[,1:2], col=diagnosis)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q18. Which of these new patients should we prioritize for follow up based on your results?

Patient 2.