

Class15 RNASeq Analysis

Gabrielle Meza (A13747395)

11/17/2021

#Background

Today we're examining a published RNA-seq experiment where airway smooth muscle cells were treated with dexamethasone, a synthetic glucocorticoid steroid with anti-inflammatory effects (Himes et al.)

We need: 1) Count of Data 2) col of data

The count of data, is the actual data, with each gene as a column header. the col is the metadata, and each gene is a row header

```
counts <- read.csv("airway_scaledcounts.csv", stringsAsFactors = FALSE, row.names=1)
metadata <- read.csv("airway_metadata.csv", stringsAsFactors = FALSE)
```

```
head(counts)
```

```
##                SRR1039508 SRR1039509 SRR1039512 SRR1039513 SRR1039516
## ENSG00000000003          723          486          904          445          1170
## ENSG00000000005           0           0           0           0           0
## ENSG00000000419          467          523          616          371          582
## ENSG00000000457          347          258          364          237          318
## ENSG00000000460           96           81           73           66          118
## ENSG00000000938           0           0           1           0           2
##                SRR1039517 SRR1039520 SRR1039521
## ENSG00000000003          1097          806          604
## ENSG00000000005           0           0           0
## ENSG00000000419          781          417          509
## ENSG00000000457          447          330          324
## ENSG00000000460           94          102           74
## ENSG00000000938           0           0           0
```

```
head(metadata)
```

```
##          id    dex celltype    geo_id
## 1 SRR1039508 control   N61311 GSM1275862
## 2 SRR1039509 treated   N61311 GSM1275863
## 3 SRR1039512 control   N052611 GSM1275866
## 4 SRR1039513 treated   N052611 GSM1275867
## 5 SRR1039516 control   N080611 GSM1275870
## 6 SRR1039517 treated   N080611 GSM1275871
```

Q1. How many genes are in this dataset?

38694 total genes

```
nrow(counts)
```

```
## [1] 38694
```

Q2. How many 'control' cell lines do we have?

4

```
sum(metadata$dex == "control")
```

```
## [1] 4
```

Side note:

lets check the correspondence of the metadata and count data setup

```
metadata $id
```

```
## [1] "SRR1039508" "SRR1039509" "SRR1039512" "SRR1039513" "SRR1039516"  
## [6] "SRR1039517" "SRR1039520" "SRR1039521"
```

```
colnames(counts)
```

```
## [1] "SRR1039508" "SRR1039509" "SRR1039512" "SRR1039513" "SRR1039516"  
## [6] "SRR1039517" "SRR1039520" "SRR1039521"
```

We can use the == thing to see if they columns and rows are the same

```
metadata$id == colnames(counts)
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

We are going to wrap this in an all thing so it will tell us if they all are true or not. if there is a false anywhere, then it would not give out TRUE. We would use this for a larger dataset potentially.

```
all(metadata$id == colnames(counts))
```

```
## [1] TRUE
```

Compare control to treated

first we need to access all the control columns in our counts data.

```
metadata$dex == "control"
```

```
## [1] TRUE FALSE TRUE FALSE TRUE FALSE TRUE FALSE
```

```
control.inds <- metadata$dex == "control"
```

```
#this is pulling out the rows that are control in the metadata sheet
metadata[ control.inds, ]
```

```
##           id      dex celltype    geo_id
## 1 SRR1039508 control   N61311 GSM1275862
## 3 SRR1039512 control   N052611 GSM1275866
## 5 SRR1039516 control   N080611 GSM1275870
## 7 SRR1039520 control   N061011 GSM1275874
```

```
#now use $id to get the ids that are controls and make it a new value
```

```
control.ids <- metadata[ control.inds, ]$id
```

use these ids to access just the control columns of our counts data. Use head because it would be a large dataset

```
head(counts[ , control.ids])
```

```
##           SRR1039508 SRR1039512 SRR1039516 SRR1039520
## ENSG000000000003      723        904        1170        806
## ENSG000000000005         0         0         0         0
## ENSG000000000419      467        616        582        417
## ENSG000000000457      347        364        318        330
## ENSG000000000460       96         73        118        102
## ENSG000000000938         0          1          2          0
```

```
control.mean <- rowMeans(counts[ , control.ids])
head(control.mean)
```

```
## ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
##           900.75           0.00           520.50           339.75           97.25
## ENSG000000000938
##           0.75
```

Q4. Follow the same procedure for the treated samples

```
treated.inds <- metadata$dex == "treated"
treated.ids <- metadata[ treated.inds, ]$id
treated.mean <- rowMeans(counts[ , treated.ids])
head(treated.mean)
```

```
## ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
##           658.00           0.00           546.00           316.50           78.75
## ENSG000000000938
##           0.00
```

we will combine our meancount data for bookkeeping purposes

```
meancounts <- data.frame(control.mean, treated.mean)
```


there are 38694 rows/genes in this dataset. You can do this in the printout and have it run the code in the printout with ** but you would have to reformat in the **

Compare the control and treated

Q5. Create a scatter plot showing the mean of the treated samples against the mean of the control samples.

Quick plot of our progress so far

```
plot(meancounts)
```




there should be many more points shown (38,000!) this could better be represented by a log scale

```
plot(meancounts, log="xy")
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted
## from logarithmic plot
```

```
## Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted
## from logarithmic plot
```



we often use log transformations as they make life much nicer in this world...

```
meancounts$log2fc <- log2(meancounts[, "treated.mean"] / meancounts[, "control.mean"])
head(meancounts)
```

##	control.mean	treated.mean	log2fc
## ENSG000000000003	900.75	658.00	-0.45303916
## ENSG000000000005	0.00	0.00	NaN
## ENSG000000000419	520.50	546.00	0.06900279
## ENSG000000000457	339.75	316.50	-0.10226805
## ENSG000000000460	97.25	78.75	-0.30441833
## ENSG000000000938	0.75	0.00	-Inf

We need to drop the zero count genes/ rows!

```
head(meancounts[, 1:2] == 0)
```

##	control.mean	treated.mean
## ENSG000000000003	FALSE	FALSE
## ENSG000000000005	TRUE	TRUE
## ENSG000000000419	FALSE	FALSE
## ENSG000000000457	FALSE	FALSE
## ENSG000000000460	FALSE	FALSE
## ENSG000000000938	FALSE	TRUE

The `which()` function tells us the indices of TRUE entries in a logical vector.

```
which (c(T,F,T))
```

```
## [1] 1 3
```

However, it is not that useful in default mode on our type of multi- column input

```
inds <- which(meancounts[,1:2] == 0, arr.ind=TRUE)
head(inds)
```

```
##           row col
## ENSG000000000005    2  1
## ENSG000000004848   65  1
## ENSG000000004948   70  1
## ENSG000000005001   73  1
## ENSG000000006059  121  1
## ENSG000000006071  123  1
```

I only care about the rows here)if there is a zero in any column I will exclude this row eventually). You then want to use `unique` to get the rows you want to look at

```
to.rm <- unique(inds[,1])
mycounts <- (meancounts[-to.rm,])
```

```
head(meancounts[-to.rm,])
```

```
##           control.mean treated.mean      log2fc
## ENSG000000000003      900.75      658.00 -0.45303916
## ENSG000000000419      520.50      546.00  0.06900279
## ENSG000000000457      339.75      316.50 -0.10226805
## ENSG000000000460       97.25       78.75 -0.30441833
## ENSG000000000971     5219.00     6687.50  0.35769358
## ENSG000000001036     2327.00     1785.75 -0.38194109
```

we now have 21817 genes remaining.

```
nrow(mycounts)
```

```
## [1] 21817
```

How many of these genes are upregulated at the log 2 fold change threshold of +2 or greater?

```
sum(mycounts$log2fc > +2)
```

```
## [1] 250
```

what percentage is this?

```
round((sum(mycounts$log2fc > +2) / nrow(mycounts))*100, 2)
```

```
## [1] 1.15
```

How about the downregulated genes?

```
sum(mycounts < -2)
```

```
## [1] 367
```

DESeq2 analysis

```
library(DESeq2)
```

```
## Loading required package: S4Vectors
```

```
## Loading required package: stats4
```

```
## Loading required package: BiocGenerics
```

```
## Loading required package: parallel
```

```
##
```

```
## Attaching package: 'BiocGenerics'
```

```
## The following objects are masked from 'package:parallel':
```

```
##
```

```
##   clusterApply, clusterApplyLB, clusterCall, clusterEvalQ,  
##   clusterExport, clusterMap, parApply, parCapply, parLapply,  
##   parLapplyLB, parRapply, parSapply, parSapplyLB
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##   IQR, mad, sd, var, xtabs
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##   anyDuplicated, append, as.data.frame, basename, cbind, colnames,  
##   dirname, do.call, duplicated, eval, evalq, Filter, Find, get, grep,  
##   grepl, intersect, is.unsorted, lapply, Map, mapply, match, mget,  
##   order, paste, pmax, pmax.int, pmin, pmin.int, Position, rank,  
##   rbind, Reduce, rownames, sapply, setdiff, sort, table, tapply,  
##   union, unique, unsplit, which.max, which.min
```

```
##
```

```
## Attaching package: 'S4Vectors'
```

```

## The following objects are masked from 'package:base':
##
##     expand.grid, I, unname

## Loading required package: IRanges

## Loading required package: GenomicRanges

## Loading required package: GenomeInfoDb

## Loading required package: SummarizedExperiment

## Loading required package: MatrixGenerics

## Loading required package: matrixStats

##
## Attaching package: 'MatrixGenerics'

## The following objects are masked from 'package:matrixStats':
##
##     colAlls, colAnyNAs, colAnys, colAvgsPerRowSet, colCollapse,
##     colCounts, colCummaxs, colCummins, colCumprods, colCumsums,
##     colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,
##     colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,
##     colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,
##     colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,
##     colWeightedMeans, colWeightedMedians, colWeightedSds,
##     colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgsPerColSet,
##     rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,
##     rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,
##     rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,
##     rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,
##     rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,
##     rowWeightedMads, rowWeightedMeans, rowWeightedMedians,
##     rowWeightedSds, rowWeightedVars

## Loading required package: Biobase

## Welcome to Bioconductor
##
##     Vignettes contain introductory material; view with
##     'browseVignettes()'. To cite Bioconductor, see
##     'citation("Biobase")', and for packages 'citation("pkgname)".

##
## Attaching package: 'Biobase'

## The following object is masked from 'package:MatrixGenerics':
##
##     rowMedians

```



```
## The following objects are masked from 'package:matrixStats':
##
##     anyMissing, rowMedians
```

We first need to setup the DESeq

design is where in the col data do we care about?

```
dds <- DESeqDataSetFromMatrix(countData=counts,
                              colData=metadata,
                              design=~dex)
```

```
## converting counts to integer mode
```

```
## Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
## design formula are characters, converting to factors
```

Run the DESeq analysis pipeline. it is doing a lot of other analysis, like p-value and stuff. it is being saved as a dds with all of the results. you can use the package default `results()` to read the data and out it into a new variable to then to pull it up and read it all.

```
dds <- DESeq(dds)
```

```
## estimating size factors
```

```
## estimating dispersions
```

```
## gene-wise dispersion estimates
```

```
## mean-dispersion relationship
```

```
## final dispersion estimates
```

```
## fitting model and testing
```

```
res <- results(dds)
head(res)
```

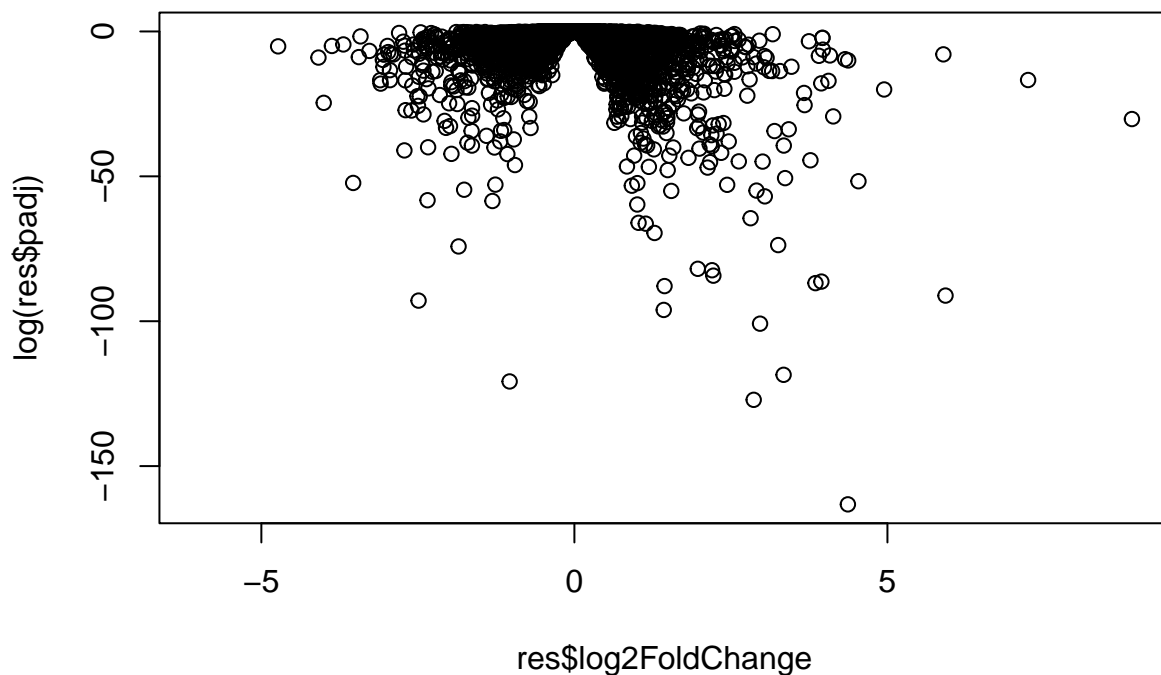
```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## Dataframe with 6 rows and 6 columns
##           baseMean log2FoldChange      lfcSE      stat      pvalue
##           <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003 747.194195      -0.3507030 0.168246 -2.084470 0.0371175
## ENSG000000000005  0.000000           NA           NA           NA           NA
## ENSG000000000419 520.134160      0.2061078 0.101059  2.039475 0.0414026
## ENSG000000000457 322.664844      0.0245269 0.145145  0.168982 0.8658106
## ENSG000000000460  87.682625     -0.1471420 0.257007 -0.572521 0.5669691
## ENSG000000000938  0.319167     -1.7322890 3.493601 -0.495846 0.6200029
##           padj
##           <numeric>
```

```
## ENSG000000000003 0.163035
## ENSG000000000005 NA
## ENSG000000000419 0.176032
## ENSG000000000457 0.961694
## ENSG000000000460 0.815849
## ENSG000000000938 NA
```

A volcano plot

This is a very common data viz of this type of data that does not really look like a volcano. you can use base r package to plot this

```
plot(res$log2FoldChange, log(res$padj))
```




when you take a log, the values you are interested in at the very low numbers since it is a p-value. To make this easier visually, you can put a negative sign in front of log

with log, as you move up in expression, go more positive. as you go down in expression, go more negative

use abline to add a line at certain points. here we are making the line at the set pValue (horizontal) to be significant. And then the cut off if there is a 2 fold change in either neg or positive direction.

```
plot(res$log2FoldChange, -log(res$padj))
abline(v=c(-2,2), col="gray")
abline(h=-log(0.05), col="gray")
```



Adding annotation data

Now lets color and annotate this data further! we want to add meaningful gene names to our dataset so we can make some sense of what is going on here!

For this we will be installing and using 2 more packages in biocManager. I installed these in the console. **AnnotationDbi** the other contains data we are going to map between and is **org.Hs.eg.db**

```
library("AnnotationDbi")
library("org.Hs.eg.db")
```

```
##
```

```
columns(org.Hs.eg.db)
```

```
## [1] "ACCNUM"      "ALIAS"       "ENSEMBL"     "ENSEMBLPROT" "ENSEMBLTRANS"
## [6] "ENTREZID"    "ENZYME"      "EVIDENCE"    "EVIDENCEALL"  "GENENAME"
## [11] "GENETYPE"    "GO"          "GOALL"       "IPI"          "MAP"
## [16] "OMIM"        "ONTOLOGY"    "ONTOLOGYALL" "PATH"         "PFAM"
## [21] "PMID"        "PROSITE"     "REFSEQ"      "SYMBOL"       "UCSCKG"
## [26] "UNIPROT"
```

Here we map to "SYMBOL" the comon gene name that the world understands and wants

```
res$symbol <- mapIds(org.Hs.eg.db,
                    keys=row.names(res),
                    keytype="ENSEMBL",
                    column="SYMBOL",
                    multiVals="first")
```

```
## 'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 7 columns
##
```

	baseMean	log2FoldChange	lfcSE	stat	pvalue
##	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
## ENSG000000000003	747.194195	-0.3507030	0.168246	-2.084470	0.0371175
## ENSG000000000005	0.000000	NA	NA	NA	NA
## ENSG0000000000419	520.134160	0.2061078	0.101059	2.039475	0.0414026
## ENSG0000000000457	322.664844	0.0245269	0.145145	0.168982	0.8658106
## ENSG0000000000460	87.682625	-0.1471420	0.257007	-0.572521	0.5669691
## ENSG0000000000938	0.319167	-1.7322890	3.493601	-0.495846	0.6200029

```
##
```

	padj	symbol
##	<numeric>	<character>
## ENSG0000000000003	0.163035	TSPAN6
## ENSG0000000000005	NA	TNMD
## ENSG0000000000419	0.176032	DPM1
## ENSG0000000000457	0.961694	SCYL3
## ENSG0000000000460	0.815849	C1orf112
## ENSG0000000000938	NA	FGR

Let's now save this data

Let's finally save our results as a .csv file so we can use for a later time, and use in different area if needed.

```
write.csv(res, file = "allmyresults.csv")
```

DAY 2- Pathway analysis

Let's try to bring some biology insight back into this work. for this we will use KEGG. So let's load up KEGG for pathway analysis.

For this, we will have to install some stuff `BiocManager::install(c("pathview", "gage", "gageData"))`

```
library(pathview)
```

```
## #####
## Pathview is an open source software package distributed under GNU General
## Public License version 3 (GPLv3). Details of GPLv3 is available at
## http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to
```

```
## formally cite the original Pathview paper (not just mention it) in publications
## or products. For details, do citation("pathview") within R.
##
## The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG
## license agreement (details at http://www.kegg.jp/kegg/legal.html).
## #####
```

```
library(gage)
```

```
##
```

```
library(gageData)
```

```
data(kegg.sets.hs)
```

```
head(kegg.sets.hs, 2)
```

```
## $'hsa00232 Caffeine metabolism'
## [1] "10" "1544" "1548" "1549" "1553" "7498" "9"
##
## $'hsa00983 Drug metabolism - other enzymes'
## [1] "10" "1066" "10720" "10941" "151531" "1548" "1549" "1551"
## [9] "1553" "1576" "1577" "1806" "1807" "1890" "221223" "2990"
## [17] "3251" "3614" "3615" "3704" "51733" "54490" "54575" "54576"
## [25] "54577" "54578" "54579" "54600" "54657" "54658" "54659" "54963"
## [33] "574537" "64816" "7083" "7084" "7172" "7363" "7364" "7365"
## [41] "7366" "7367" "7371" "7372" "7378" "7498" "79799" "83549"
## [49] "8824" "8833" "9" "978"
```

Before we can use KEGG we need to get our gene identifiers in the correct format for KEGG, which is ENTREZ format in this case. So let's look into the dataset again to get familiar with the language

```
head(rownames(res))
```

```
## [1] "ENSG00000000003" "ENSG00000000005" "ENSG000000000419" "ENSG000000000457"
## [5] "ENSG000000000460" "ENSG000000000938"
```

```
columns(org.Hs.eg.db)
```

```
## [1] "ACCNUM" "ALIAS" "ENSEMBL" "ENSEMBLPROT" "ENSEMBLTRANS"
## [6] "ENTREZID" "ENZYME" "EVIDENCE" "EVIDENCEALL" "GENENAME"
## [11] "GENETYPE" "GO" "GOALL" "IPI" "MAP"
## [16] "OMIM" "ONTOLOGY" "ONTOLOGYALL" "PATH" "PFAM"
## [21] "PMID" "PROSITE" "REFSEQ" "SYMBOL" "UCSCKG"
## [26] "UNIPROT"
```

The main **gage()** function requires a named vector of fold changes, where the names of the values are the Entrez gene IDs.

Note that we used the `mapIDs()` function above to obtain Entrez gene IDs (stored in `resentrez`) and we have the fold change results

```
res$entrez <- mapIds(org.Hs.eg.db,
                    keys=row.names(res),
                    keytype="ENSEMBL",
                    column="ENTREZID",
                    multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
res$genename <- mapIds(org.Hs.eg.db,
                      keys=row.names(res),
                      keytype="ENSEMBL",
                      column="GENENAME",
                      multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

Assign names to this vector that are the gene IDs that KEGG wants

```
foldchanges = res$log2FoldChange
names(foldchanges) = res$entrez
head(foldchanges)
```

```
##          7105          64102          8813          57147          55732          2268
## -0.35070302          NA  0.20610777  0.02452695 -0.14714205 -1.73228897
```

There we go! now let's run the **gage()** pathway analysis. Get the results by doing `keggres =` and the `gage` function, and it saves the data as `keggres`

We can then look at the `attributes()` of this or indeed any R object

```
keggres = gage(foldchanges, gsets=kegg.sets.hs)
attributes(keggres)
```

```
## $names
## [1] "greater" "less"    "stats"
```

```
head(keggres$greater)
```

```
##                                p.geomean stat.mean      p.val
## hsa00500 Starch and sucrose metabolism  0.002822007  2.825461 0.002822007
## hsa00330 Arginine and proline metabolism 0.012317455  2.280002 0.012317455
## hsa04910 Insulin signaling pathway      0.017110962  2.129511 0.017110962
## hsa04510 Focal adhesion                 0.025239833  1.961955 0.025239833
## hsa04920 Adipocytokine signaling pathway 0.043426078  1.725063 0.043426078
## hsa00790 Folate biosynthesis            0.048254489  1.744387 0.048254489
##                                q.val set.size      exp1
## hsa00500 Starch and sucrose metabolism  0.6010875      54 0.002822007
## hsa00330 Arginine and proline metabolism 0.7774866      54 0.012317455
## hsa04910 Insulin signaling pathway      0.7774866     138 0.017110962
## hsa04510 Focal adhesion                 0.7774866     200 0.025239833
## hsa04920 Adipocytokine signaling pathway 0.7774866      68 0.043426078
## hsa00790 Folate biosynthesis            0.7774866      11 0.048254489
```

```
head(res)
```

```
## log2 fold change (MLE): dex treated vs control
## Wald test p-value: dex treated vs control
## DataFrame with 6 rows and 9 columns
##           baseMean log2FoldChange      lfcSE      stat      pvalue
##           <numeric>      <numeric> <numeric> <numeric> <numeric>
## ENSG000000000003 747.194195      -0.3507030  0.168246 -2.084470 0.0371175
## ENSG000000000005   0.000000          NA          NA          NA          NA
## ENSG000000000419 520.134160      0.2061078  0.101059  2.039475 0.0414026
## ENSG000000000457 322.664844      0.0245269  0.145145  0.168982 0.8658106
## ENSG000000000460  87.682625     -0.1471420  0.257007 -0.572521 0.5669691
## ENSG000000000938  0.319167     -1.7322890  3.493601 -0.495846 0.6200029
##           padj      symbol      entrez      genename
##           <numeric> <character> <character>      <character>
## ENSG000000000003  0.163035      TSPAN6       7105      tetraspanin 6
## ENSG000000000005          NA      TNMD       64102      tenomodulin
## ENSG000000000419  0.176032      DPM1       8813      dolichyl-phosphate m..
## ENSG000000000457  0.961694      SCYL3      57147      SCY1 like pseudokina..
## ENSG000000000460  0.815849      C1orf112     55732      chromosome 1 open re..
## ENSG000000000938          NA      FGR       2268      FGR proto-oncogene, ..
```

The pathview() function will add our genes to a KEGG pathway as colored entries:

```
pathview(gene.data=foldchanges, pathway.id="hsa05310")
```

```
## 'select()' returned 1:1 mapping between keys and columns
```

```
## Info: Working in directory /Users/Gabriellemeza/Desktop/UCSD/Bioinformatics/bggn213_github/Class15
```

```
## Info: Writing image file hsa05310.pathview.png
```